

Generative AI

with Prof. Anne-Marie Kermarrec, EPFL & Sasha Rubel, AWS
Board Agora Discussion Session / 23rd April 2026 / Paris

Executive Summary

European boards are caught between two clocks running at fundamentally different speeds. The clock of AI innovation is moving at an unprecedented pace, the shift from generative to agentic AI took less than a year, a transformation that required a decade in previous technology cycles. The clock of governance, regulatory approval, and board oversight moves far more slowly.

This session established three core findings:

- The majority of European companies are not yet formally overseeing AI at board level in the way that regulators will require. The gap between deployment and governance is widening, not closing.
- Closing that gap is not principally a technical challenge. It is a cultural and leadership challenge, requiring boards to rewire their mental models, build new muscles, and embrace distributed accountability rather than concentrated control.
- The starting point is not an AI strategy. It is data governance. Boards that skip this foundational work will find their AI investments built on unstable ground.

"Today, six out of ten large European companies are formally not overseeing AI at board level the way they should – and the way regulators will require from August on."

I. THREE MYTHS BOARDS MUST DISPEL

Effective board governance of AI begins with an accurate picture of what AI is and is not. Three persistent myths distort that picture in ways that are consequential for how boards ask questions, set risk appetite, and hold management accountable.

Myth 1: AI is a black box we cannot understand

AI models are complex, and even their designers cannot fully explain the role of every individual parameter. This is true. But it does not mean boards are helpless. The more important questions, what a specific system has been optimised for, what data was used in training, where it has been tested and validated, are entirely answerable. Boards do not need to understand the mathematics. They need to be informed beyond what their CTO tells them. That is an achievable standard.

Myth 2: If it works, it is safe

A model that is accurate 90% of the time may nonetheless be failing in exactly the cases that matter most, the most vulnerable customers, the highest-stakes decisions, the edge cases that define fiduciary exposure. Accuracy is not safety. The question is not whether the model performs on average, but whether it performs where it counts, and what happens when it does not.

The bias embedded in training data is one of the clearest illustrations. A model trained on systematically skewed data produces systematically skewed outputs, even if no one intended it, and even if headline performance metrics look strong. A system designed to address the side effects of menopause, trained on a dataset that was 90% male health data, will produce unreliable results regardless of its accuracy score. This kind of structural mismatch between data and purpose is not an edge case. It is common. And it is precisely the kind of failure that boards are positioned to prevent, by asking the right questions before deployment, not after harm occurs.

Myth 3: AI can do everything or soon will

AI can perform certain extraordinarily complex tasks, modelling protein structures at a level beyond human capability, while failing at tasks a six-year-old could manage, such as predicting the trajectory of a ball from video. Boards that overestimate AI capability take on unmanaged risk. Boards that underestimate it miss strategic opportunity. The goal is calibrated understanding: knowing which parts of the business can genuinely benefit from which tools, rather than deploying AI because it is expected, or refusing to deploy it because it feels opaque.

AI has moved from circus animal – doing impressive tricks – to Labrador: always there, occasionally making a mess, but impossible to ignore. The danger with the Labrador is that familiarity breeds complacency, and complacency stops boards from asking the most fundamental questions about risk.

II. THE GOVERNANCE GAP IS REAL – AND GROWING

The session opened with a striking framing: the majority of large European companies approve AI strategy but do not formally oversee AI governance. There is a critical distinction between approving how a company plans to deploy AI and actively overseeing how AI is being managed and governed once deployed. Boards are doing the former. Most are not yet doing the latter.

Experience across financial services, healthcare, and technology confirmed this picture. The challenge is less about awareness than about pace. Boards understand that AI matters. What they struggle to keep up with is the velocity at which the landscape is changing.

In major European banks, development cycles have compressed from six months to two days for certain product categories. The control infrastructure, risk assessment processes, and regulatory approval timelines were built for the old pace. They are structurally misaligned with the new one. A solution that takes two days to build now

takes sixteen months to obtain regulatory approval and by the time that approval arrives, the solution may already be obsolete.

The move from dial-up to 3G took ten years. The move from generative AI to agentic AI took less than one. Governance frameworks cannot be written for a technology that has already moved on.

This is not an isolated case. It is the structural reality facing boards across every sector where AI is being deployed at scale. The tension between innovation velocity and regulatory process has no easy resolution – but it demands a governance response. Boards that wait for the right regulatory moment to act will find, when the moment arrives, that they are already behind.

III. LEADERSHIP IN THE AGE OF AI: THREE SHIFTS

The governance challenge is, at its core, a leadership challenge. The technological revolution is simultaneously a cultural revolution, and the cultural dimension is where boards are least prepared. Three shifts in leadership posture emerged as central to the discussion.

From having the answers to asking the right questions

Traditional board leadership presumed that seniority confers answers. In a landscape where even leading researchers describe their field as changing faster than they can track, with more than 400 academic papers on AI published daily, no individual holds the answers. The value of board leadership has shifted toward identifying the right questions: about risk, about accountability, about what responsible deployment means in practice for a given organisation.

The cultural barrier is real. The shame of not knowing, the instinct to close the door before asking a basic question about how a technology works, costs organisations dearly. A minister responsible for digital transformation who cannot admit to not understanding the tools he is meant to govern is not an anomaly. He is a pattern. The sign of intelligence in this environment is precisely the willingness to say: I do not understand this yet. Teach me.

From control to curiosity

Risk-averse cultures, and European boardrooms skew risk-averse, default to control as the primary governance response. But control, in a fast-moving environment, can become paralysis. The board that waits for a competitor to fail before it acts will find, by the time it acts, that the competitor has moved two steps further ahead. The old playbook, let the frontrunner fail, then copy, no longer works when the frontrunner's next product is already live before the copy is approved.

Curiosity is not a softening of governance discipline. It is a prerequisite for exercising it effectively. AWS's first-day expectation, that if you do not fail hard in your first year, you have not stretched far enough, is the cultural counterpoint to everything boards have been trained to believe about risk. Embracing the possibility of failure,

learning from it, and correcting course quickly is the only governance model that works at the pace AI is moving.

From centralised to distributed accountability

The most common governance instinct when faced with a new risk is to appoint someone to own it. The Chief Responsible AI Officer. The AI Risk Committee. The designated board member. These structures are not wrong, they are insufficient. AI touches every function, every workflow, every customer relationship. Responsibility must be distributed across the organisation, not delegated to a single point. But distribution without clarity creates diffusion. The governance challenge is to make everyone accountable while being absolutely clear about what each level of the organisation is accountable for and to ensure that accountability reaches the board level, not just stops at the CTO's desk.

Old Leadership Model	Required Leadership Model
I have the answers	I know what questions to ask
Control	Curiosity
Centralised accountability	Distributed accountability
Responsibility bolted on	Responsibility built in by design
AI strategy first	Data governance first

IV. THE REGULATORY BIND AND WHAT BOARDS CAN DO ABOUT IT

The EU AI Act emerged as both a structural reference point and a source of significant operational anxiety. The numbers are stark:

50% +	of jobs in the next three years will require digital skills – cited as one of the primary blockers to AI adoption across European businesses
--------------	--

42%	of overall IT budget in the EU goes to compliance — rising to 47% in France
------------	---

68%	of EU businesses do not understand their obligations under the EU AI Act
------------	--

4 in 10	startups would consider leaving Europe due to regulatory complexity; for advanced organisations, that rises to more than half
----------------	---

41%	of European businesses cite regulatory fragmentation within the EU as a major barrier to scale
------------	--

When the rules are so opaque that every company is perpetually uncertain about whether it is in breach, regulation stops functioning as a governance framework and starts functioning as a tool of selective enforcement. The concern – heard clearly in the room – is that AI regulation with disproportionate, ill-defined fines tied to transaction volumes or turnover creates exactly that dynamic: a system where the rational response is concealment rather than compliance, and where development migrates to jurisdictions with clearer, lighter frameworks.

A healthcare startup competing globally cannot afford to spend 47% of its IT budget on compliance when a competitor in Brazil or Japan spends 22%. The compliance spend is not the objection. The opacity is. Startups in Spain that would love to expand to France and Germany stay small because the two countries are interpreting the EU AI Act differently and no one wants to take the risk of being wrong.

The history of the Act itself illustrates the structural challenge. Drafting began in 2018. By the time it was adopted, generative AI had arrived – technology that did not exist when the drafting started. A new title and chapter had to be added to address general-purpose AI. The question this raises is not abstract: how do you build regulatory frameworks for technology that has not been invented yet, without either under-regulating what already exists or over-constraining what is still to come?

The call from the room was not for deregulation. It was for clarity. Boards have a constructive role to play here – not as lobbyists, but as informed participants who can bridge the gap between business reality and policy process. The disconnect between what regulators design and what companies can practically implement is, in part, a consequence of too little dialogue. Boards that understand both worlds are well-placed to help close it.

"Hallucinations in AI systems are not a bug. They are a feature. The question is not how to eliminate them – it is how to design human oversight that accounts for them."

V. HOW BOARDS ARE RESPONDING: GOVERNANCE MODELS IN PRACTICE

The most substantive exchange in the session moved from diagnosis to practice: what boards are actually doing to close the governance gap, rather than what they aspire to do.

The AI Advisory Board

After a transparent capability review – the results of which are published on the bank's website – one of Europe's leading banks concluded that no board member was currently capable of meaningfully challenging executive management on AI. The first instinct was to recruit an expert non-executive director. Every suitable candidate was in an active operational role; joining a bank board – time-consuming, heavily regulated – was not an attractive proposition.

The solution was an AI advisory board, deliberately structured outside the formal board architecture but directly connected to it. Its composition will include two board members (including the chairman), the CEO, the head of private wealth management, and three to five external members drawn from three talent profiles: leaders who have transformed a legacy company into an AI-driven one; founders who built data and technology businesses from scratch; and practitioners with deep engineering and technical experience.

Twelve candidates across three countries have been identified. What proved surprising: some of the most senior AI leaders from major competing European banks expressed interest in joining – not despite being busy, but because the format offered something rare: structured peer dialogue across industries on a challenge everyone is navigating in isolation. The advisory board will meet three times a year. Its function is not to approve strategy – that remains with the board – but to bring genuine operational challenge to executive management's thinking, and to help board members calibrate the quality of the questions they are asking.

The Ethics and Data Committee

Orange developed a committee on data and AI ethics that assembled individuals from diverse professional and disciplinary backgrounds into an ongoing advisory function. The emphasis on breadth – professional, disciplinary, and demographic diversity – was a deliberate design choice. The argument: responsible AI governance requires as many different vantage points as the population that will ultimately use the technology. A committee of technologists and lawyers will miss what a committee that also includes sociologists, healthcare practitioners, and communications experts will catch.

The Challenge Culture

A third approach focused not on structure but on habit. A board with a background in the energy sector had developed a practice of systematically stress-testing the long-term strategic plan from unexpected angles: not only 'is this plan correct?' but 'what are the scenarios in which this plan fails, and how quickly could we adapt?' Challenging the plan on solar deployment assumptions, on battery cost trajectories, on geopolitical discontinuities – this built a collective language for uncertainty and a governance culture less dependent on any single expert. The mental muscle it develops – the ability to hold and interrogate multiple futures simultaneously – transfers directly to AI governance.

"There is a mental model within boards: risk oversight, concentrated responsibility, long-term planning. AI doesn't fit that model. Boards that try to govern AI with the tools they built for a slower world will find those tools don't work."

VI. THE SKILLS BOARDS AND COMPANIES ACTUALLY NEED

The question of skills permeated the discussion at every level: in the boardroom, across the workforce, and in the next generation of leaders. It is also where the gap between the technical narrative and the human reality is most visible.

What businesses say they want

The number one question government officials ask behind closed doors – far more often than questions about regulation – is: what should my daughter study so she has a job tomorrow? The answer offered in this session was: English literature. Not as a provocation, but as shorthand for a specific capability set that businesses, when surveyed, consistently rank as more valuable than narrowly technical skills.

What businesses say they want from the workforce of the future is: the ability to translate business requirements into technical specifications; systems thinking that connects apparently unrelated functions and silos; the capacity to redesign and reinvent processes from scratch; creative and critical thinking; and the ability to collaborate – with people and with AI tools – in ways that amplify rather than replace human judgment.

50% +

of jobs in the next three years will require digital skills – cited as one of the primary blockers to AI adoption across European businesses

The risk is that organisations focus upskilling effort too narrowly on technical capacity. Not everyone will become an AI developer. Everyone will use AI. The skills needed to use it well – to ask good questions, evaluate outputs critically, understand the stakes of bias and privacy, and know when human judgment must override the model – are the soft skills that sit alongside, not beneath, technical literacy. Boards that upskill their organisations purely on the technical dimension are solving half the problem.

What boards specifically need

At board level, the challenge carries an additional constraint: experience in digital transformation from five years ago does not translate automatically into AI governance competence today. The pace of change is too great, and the nature of the risk is different. The goal is not to make every board member an AI expert. It is to build a board that collectively holds the right combination of experience – and to supplement that with advisory structures providing the technical depth the board itself cannot credibly hold.

The data literacy problem

Before organisations can use AI effectively, they need to become genuinely data driven. Finance functions tend to be data-disciplined by culture and necessity. Marketing functions, historically, have been less so. Getting every part of an organisation to treat data as an asset requiring active stewardship – one dataset, one owner, clear accountability – is a change management challenge that sits entirely upstream of any AI deployment. Boards have a direct role here: data governance is not a technical decision. It is an organisational one, and it requires the kind of cross-functional authority that only the board can provide.

"Garbage in, garbage out. But before you can fix what goes in, you need structures that create accountability for it – and that is a board-level conversation, not a CTO conversation."

VII. RESPONSIBLE AI – FROM PRINCIPLE TO PRACTICE

The term 'responsible AI' was interrogated with some scepticism. Most policymakers, asked whether principles like transparency, privacy, robustness, and controllability matter, say yes without hesitation. Most engineers, asked to build systems that embody those principles, respond that they are well-intentioned abstractions with no operational meaning. The gap between aspiration and practice is where responsible AI frameworks must do their work.

Responsibility by design, not as afterthought

The most effective responsible AI frameworks embed responsibility into the development process from the outset – not as a compliance layer applied after the fact, but as a design constraint that shapes what gets built and how. Constitutional AI approaches, which encode human rights principles directly into AI system architecture, represent the direction of travel: responsibility that is native to the system rather than bolted onto it.

The business case for this approach is not only ethical. Research suggests that a responsible-by-design approach can double profit margins. The reasoning is straightforward: people consume technology they trust. Trust is earned through consistency, transparency, and alignment of values. Companies that build trust into their AI systems – rather than treating it as a reputational add-on – create a durable competitive advantage. Recent events involving model providers that made decisions misaligned with their users' values demonstrated, concretely, that ethics and commercial interest are not as separable as some governance frameworks assume.

Diversity as a design input, not a compliance requirement

Responsible AI is not only about what is built, but who is in the room when it is built. The training data examples that surfaced in the session – a gossip-algorithm illustration generating images of blonde women, a menopause solution trained on male health data – are symptoms of a deeper structural problem: the people designing and deploying AI systems do not yet reflect the diversity of the populations those systems will affect. This is not a soft concern. It is a governance failure with concrete consequences: biased outputs, regulatory exposure, and the erosion of trust among precisely the customers the technology was meant to serve.

Macroeconomic research shows that boards with more than one woman produce higher return on investment and innovation. More diverse teams generate more ROI. The direction of travel matters: in France, at the current rate, gender parity is 135 years away. Against a geopolitical backdrop that is actively eroding diversity commitments, boards that treat diversity as discretionary rather than structural are making a business decision, not just an ethical one.

The human oversight question

When AI moves from decision support to autonomous action – executing transactions, managing processes, interacting with customers without a human in the loop – what does meaningful oversight actually mean? The comparison to human judgment is instructive. Research shows that AI systems deployed for high-risk decisions, including credit assessment, can exhibit lower bias than human equivalents. But there is a critical asymmetry: a human failure can be isolated, attributed, and characterised as individual error. A systemic AI failure cannot. When the model is wrong at scale, it is a governance failure – and the regulatory and reputational consequences are correspondingly larger. Boards need to design oversight that is calibrated to this asymmetry, not simply to the headline accuracy rates.

VIII. THE DATA GOVERNANCE IMPERATIVE

The session converged, at its close, on a point that had been implicit throughout the discussion: the most important thing a board can do to prepare for effective AI governance is not to develop an AI strategy. It is to get the data house in order first. This is the unglamorous work that precedes the glamorous outcome. Data quality, data provenance, data ownership structures, consent frameworks, and security protocols are the foundations on which any AI deployment rests. Approving an AI strategy without first understanding the quality and governance of the underlying data is approving a plan built on untested assumptions. The outputs of an AI system are only as reliable as the data it was trained on – and in most large European organisations, a frank assessment of data quality, consistency, and ownership would reveal significant gaps.

More than 84% of top management and board members identify responsible AI governance as a top management issue. Fewer than 25% have a strategy in place to address it. The gap is not awareness. It is the willingness to do the structural work that makes governance operational rather than aspirational – to put in place the data ownership accountability, the cross-functional data standards, and the ongoing monitoring that any serious AI deployment requires.

The original proposal for this session was to address both AI and data governance. The group concluded, after 90 minutes, that the AI conversation alone could fill the time. The data conversation will require its own session – and the demand for it was unanimous. That is itself a governance signal: data governance is not a sub-topic of AI governance. It is its prerequisite.

"The most effective organisations deploying AI are companies that did the dirty work first – building a data governance strategy, then a responsible AI strategy. They are not the ones that jumped straight to the shiny object."

TAKEAWAYS FOR BOARD LEADERS

The session was not prescriptive. It did not offer a checklist. But a clear set of actionable orientations emerged from the conversation:

- Understand what your company is actually doing with AI – not at the level of strategy documents, but at the level of live deployment. Which systems are running? What do they decide? Who is accountable for what they produce?
- Build the questioning muscle. The right questions about risk, accountability, and data quality are more valuable than premature answers. Boards that can ask better questions get better answers from management.
- Start with data governance. Any AI strategy not grounded in a clear understanding of data quality, ownership, and provenance is a strategy built on assumptions. Do the dirty work first.
- Distribute accountability – but define it clearly. Do not create a Chief Responsible AI Officer and call it governance. Responsibility must reach every level of the organisation, with unambiguous accountability at each.
- Build advisory capacity. If your board cannot currently challenge executive management on AI, that is a solvable problem. Advisory boards, external experts, structured learning programmes – the formats exist and are proving effective.
- Embrace the culture of not knowing. Boards that treat gaps in AI knowledge as vulnerability will be outpaced. Boards that treat them as the starting point for learning are better positioned to govern effectively – and to attract the advisors who can help.
- Engage with regulators as a constructive actor. The disconnect between business reality and regulatory process is, in part, a consequence of too little dialogue. Boards with operational AI experience have something valuable to contribute to that conversation.

Acknowledgements

This knowledge paper synthesises insights from the Board Agora discussion session held on 23 April 2026. We extend our gratitude to Prof. Anne-Marie Kermarrec and Sasha Rubel for their expert contributions, and to all participating board members whose candid reflections and challenges enriched this dialogue. The quality of a Chatham House conversation depends entirely on the willingness of participants to speak with genuine honesty.