

AI & the Future of Work: Leading the Human Transition

with Nathalie Scardino, Salesforce & Jean-Christophe Deslarzes, Adecco Group
Board Agora Discussion Session | 15th June 2026 | Zurich

1. FRAMING THE CONVERSATION

On 15 June 2026, Board Agora convened in Zurich a gathering of board members and governance practitioners from across Europe, joined by a senior voice from Silicon Valley. The session, the third in Board Agora's 2026 AI programming arc, following editions in Paris on AI governance and Munich on agentic AI and the EU AI Act, asked the most consequential question facing organisations today: how do companies, and the boards that oversee them, lead the human transition through artificial intelligence?

The gathering was timely. Earlier in June, Anthropic had published an open letter calling for a pause on the most advanced AI development, acknowledging that the systems now being built may well be the last designed by humans alone, future generations of AI could design their own successors. That signal, coming from one of the frontier laboratories, sharpened the room's sense of urgency: the question before boards is no longer whether AI will reshape work, but whether organisations are moving fast enough, and with enough intentionality, to bring their people through the transition intact.

The session was introduced by a three-phase maturity model. In Phase One, AI is a tool, a sophisticated chat interface or automation layer that sits alongside existing processes without fundamentally altering them. In Phase Two, AI becomes a colleague: a persistent, responsive entity that participates as a genuine member of teams, handling end-to-end workflows and enabling new kinds of collaboration. In Phase Three, already arriving in pockets of the most advanced organisations, AI becomes a system, operating autonomously, adapting itself and taking consequential decisions within defined parameters. The governance implications of each phase differ radically. The discussion that followed examined how companies and boards are experiencing – and managing – this human transition today.

The discussion therefore focused not only on understanding the shift, but on identifying what boards must concretely do differently to govern workforce transformation in the age of AI.

2. THE SALESFORCE EXPERIENCE: BUILDING A WORKFORCE READINESS TRANSFORMATION

The session opened with a detailed account of how one of the world's leading enterprise technology companies – 80,000 people across 50 countries, with EMEA as its second-largest market – has navigated the transition. Two years ago, the company's CEO presented the leadership team with a proposition: data has been important, but the era of agents is what will define the next chapter. That insight became the founding impulse of Agentforce, the company's AI-agent platform – and set in motion what Nathalie, the company's President and Chief People Officer, described as a wholesale workforce readiness transformation.

The realisation came quickly that this was not, at its core, a technology project but a workforce transformation. Generative AI is the fastest consumer adopted technology in history, higher than the personal computer or mobile phone. The technology was moving at a speed that made every prior assumption about organisational change obsolete. But underneath the technology shift was something more fundamental: work itself was changing, and workflows were being redesigned at their core.

2.1 Learning: Top-Down and Bottom-Up

The company's first response was to establish AI fluency as a universal baseline. This was pursued simultaneously from the top and from the floor. The CEO restructured the company's flagship quarterly leadership gathering – 500 executives in person, broadcast to the entire workforce – so that attendance in the room required proof of having built an agent. Every executive entered a sandbox environment and built their first agent, with product teams sitting alongside them at every table. The session did not end with a presentation: it ended with keyboards.

This was not a one-off exercise. The company ran hackathons across every team and department, inviting employees to imagine what they would automate, redesign, or simplify if given the freedom to do so. Thousands of submissions arrived. Quarterly Agentforce learning days were held in every office, giving each employee one full day of hands-on engagement with the technology. The purpose was demystification: to show that building an agent is not the exclusive province of engineers, and that the logic of AI – data, metadata, workflow redesign – is accessible to anyone willing to engage with it.

Every employee, including the CEO, was required to become an Agentblazer, certified through Trailhead, the company's online learning platform. The gamification of learning through the company's own Trailhead platform made the process competitive and visible. By the time the company had reached the second phase of its transformation, it could credibly claim that 100% of its workforce had engaged with AI in a structured, accountable way.

This approach closely mirrors Salesforce's AI Fluency Playbook, which defines AI fluency as a measurable organisational capability built through hands-on learning, leadership participation, and structured adoption frameworks – rather than passive training alone, which Salesforce provided to our members and can be found in this group's documents section.

2.2 The 4Rs Framework

Adoption alone, however, proved insufficient. As thousands of employees began building agents, a new problem emerged: proliferation without purpose. Nobody needed twelve separate onboarding agents or seventy-five variants of a seller agent. The value was diffused and hard to measure. The company needed to redesign how they worked with agents, and engaged their people in the change. In response, the company developed what Nathalie called the 4Rs Framework, a resource to support their customers in the people transformation and framing as they communicate progress to the board.

This framework aligns with Salesforce's Workforce Innovation Playbook ("The 4Rs"), which structures workforce transformation around four coordinated levers: Redesign, Reskill, Redeploy, and Rebalance – providing a practical model for translating AI deployment into measurable workforce and business outcomes (this document is also available to you in this group's documents section).

The first R is Redesign. Every job is mapped at the task level. The company employs 2,500 unique job roles encompassing nearly 100,000 distinct tasks. The Workforce Innovation team – a new function that Nathalie described as the future of HR – works role by role to determine which of those tasks should be done by a human, which by an agent, and which should be automated entirely. This is painstaking, unsexy work, but it is where the value is released. A bad agent, as Nathalie observed, will only amplify a bad process. The redesign of the underlying job must precede or accompany the deployment of the technology.

One concrete illustration of Redesign in practice is Salesforce's deployment of Slackbot – the AI assistant embedded directly in Slack, the company's primary collaboration platform. Rather than asking employees to move to a separate AI tool, Slackbot meets employees where work already happens: synthesising meeting notes, drafting communications, surfacing relevant information, managing scheduling workflows, and reasoning through complex questions in real time. For the Workforce Innovation team, Slackbot represents the Redesign principle made tangible – cognitive tasks previously owned entirely by a person, now shared between human and agent, with the human retaining judgment and the agent compressing the time and effort required to exercise it.

The second R is Reskilling. Once jobs have been redesigned, the competencies required change, and employees must be equipped for what the new job demands.

Reskilling is continuous rather than episodic, and it operates across multiple levels of fluency: basic AI literacy for the entire workforce, deeper capabilities for Agentblazer Champions, and genuine technical depth for those building and scaling the systems.

The third R is Redeployment. The company's customer success group experienced the most dramatic early disruption going from approximately 9,000 support engineers to 5,000 over the span of about a year. Agentforce began deflecting 70% of the work these employees had previously done. The response was not to reduce headcount but to redeploy employees into growth areas of the business, such as professional services and sales.

The fourth R is Rebalancing. The company now reports to its board on a regular cycle showing the precise balance of work being done by humans versus AI, and the type of AI involved in each case – basic automation, agentic AI, or multi-system orchestration. This reporting framework gives the board visibility without requiring it to become technically expert.

“The biggest lesson I’ve learned is that we were so hyper-focused on deploying the AI. It is now not about deploying the AI. It is about what work has fundamentally been redesigned with the AI – and that is the shift we are going through.”

2.3 The Ethical and Humane Use Council

The company has a long-established ethical and humane use council that reviews responsible AI use cases, including representation from the People function, IT, Legal, Security, Operations and Finance. The council's governing question is simple: just because you can, should you? Every agent or AI application must pass through this filter before it reaches production. Resume screening agents, for example, may screen the top of the hiring funnel – the company receives two million applications per year – but no hiring decision may be made by AI alone. The human remains in the loop at the point of consequence.

The council is also the body that determines global standards, ensuring that what is deployed in San Francisco is deployed under the same guardrails in Singapore or Stuttgart. As the volume of agent-building accelerates, the company is exploring ways to augment the council process with AI systems in the flow of work to enable speed of innovation while prioritizing governance, compliance, and their trusted AI principles.

3. THE ADECCO GROUP PERSPECTIVE: DATA, PURPOSE AND THE LEADERSHIP GAP

Jean-Christophe, Chair of the Board of the Adecco Group and Chair of the Board of Constellium, offered a data-grounded counterpoint to the Salesforce experience. The Adecco Group runs surveys annually across the global working population – the most recent covered 38,000 workers in 30 countries – as well as a separate survey of 2,000 leaders who together lead 8.6 million workers. The data from these surveys tells a story of rapid acceleration in individual AI adoption, accompanied by a widening gap in organisational and leadership readiness. In contrast to Salesforce’s structured transformation approach, this data highlights how uneven and uncoordinated AI adoption remains across most organisations.

3.1 Three Key Findings from the Data

The first finding concerns the pace of adoption. In 2024, 20% of the global workforce considered AI potentially interesting. By 2025, that figure had risen to 70%. By 2026, 90% of workers report using AI in some form in their daily work. The speed of this shift is without precedent in the history of technological adoption. Yet organisational readiness has not kept pace: workers report feeling unsupported, and the gap between what leaders believe their workforce can absorb and what employees actually experience is measurable and widening. Leaders estimate that 45% of workers will be actively engaged with AI agents; employees themselves put the figure at 30%.

The second finding concerns impact measurement. Employees report saving approximately two hours per day through their use of AI. This is a significant productivity gain – but it is largely invisible in the bottom line. The reason is that AI has been integrated into individual workflows without redesigning the underlying processes. The two hours saved are used for strategic thinking, quality checking, and self-training – valuable but not yet translating into organisational output or measurable business impact. Until the process is redesigned, the productivity gain is trapped at the individual level.

The third finding is that purpose remains the central driver of retention. In an environment of rapid change and rising uncertainty, employees who feel their organisation has a strong and visible sense of purpose are retained at a rate of 99%. Where purpose is absent or unclear, that retention rate falls to roughly 50%. The AI transition, which accelerates change and amplifies uncertainty, makes the presence of organisational purpose more – not less – critical.

3.2 Leadership Imperatives

Drawing on these data points, Jean-Christophe identified three leadership imperatives for the AI era.

The first is learning to learn. This has always been a leadership virtue, but the speed of the current transition makes it non-negotiable. Leaders at every level – board, executive committee, and throughout the organisation – must build structured learning into their routine. The Adecco Group itself offers 760 learning modules, tiered by depth: basic AI literacy for all employees, an ambassador track for AI champions, and advanced modules for those building and scaling solutions.

The second imperative is measuring impact. Saving two hours per individual per day is not, in itself, a business result. Leaders must close the loop between individual efficiency gains and organisational outcomes – redeploying the unlocked capacity into genuine value-creating work and tracking the result with the same rigour applied to financial metrics.

The third imperative is empathy. As Jean-Christophe noted, survey after survey – across industries, geographies, and compensation levels – shows that career development and the feeling of being supported by leadership consistently outrank pay as drivers of engagement and retention. The best boss most people have ever had, he observed, was an empathetic leader. AI does not change that; if anything, it amplifies the need for it.

"It is not AI that will take you out of the job, it is the people who use AI who will take you out of the job."

4. VOICES FROM THE BOARDROOM: CHALLENGES, EXPERIMENTS AND EMERGING PRACTICE

The subsequent discussion brought in a richly diverse set of perspectives from participants representing financial services, industrial manufacturing, automotive suppliers, energy infrastructure, construction materials, life sciences, the recruitment industry, and governance advisory. Several distinct themes emerged.

4.1 The Governance Dilemma: Strategy Without Deep Knowledge

One of the most candid observations came from a board member in a highly regulated sector, who articulated a dilemma that resonated widely around the table. When boards define business strategy, they do so from a position of deep knowledge: they understand their markets, their customers, their competitive position. With AI strategy, that depth of understanding is absent for most boards. The result is a paralysis: boards are reluctant to define a top-down strategy for something they do

not fully understand, but they are equally unwilling to allow bottom-up adoption in regulated industries where control failure has serious consequences.

The response from participants was pragmatic: the answer is not to wait for full understanding, but to establish a minimum viable governance framework and begin testing. In the current environment, the pace of change means that the cost of inaction quickly exceeds the cost of imperfect action. Several participants noted that the regulatory environment – particularly the EU AI Act – provides a useful scaffold: it requires boards to engage with AI governance even if it does not tell them exactly what to do.

4.2 The Five-Pole Tent: Orchestrating Competing Priorities

A participant with extensive board experience across regulated industries introduced a metaphor that anchored the governance discussion for the rest of the session. Successful AI governance at the board level, she argued, requires maintaining a balance among five structural poles: technology (including AI deployment), cybersecurity, business unit ownership of P&L and outcomes, the CHRO or CPO function responsible for workforce strategy, and transformation leadership (typically the Chief Transformation Officer, a role that has proliferated since the pandemic and whose remit now increasingly centres on AI implementation).

Pull one pole too hard, and the tent collapses. An organisation that is technology-led without proportional investment in people and security becomes fragile. One that is excessively HR-driven may be slow to capture the commercial opportunity. The board's role is to act as an orchestra conductor – ensuring that all five instruments play together, that no single voice drowns out the others, and that the overall composition serves the strategic vision of the organisation.

4.3 The Revenue Disruption Question

A board advisor with direct experience across industrial conglomerates pressed the room on a question that many boards are still unwilling to ask directly: which parts of our current revenue model will become effectively free within three to five years? In businesses where significant portions of revenue derive from knowledge services, advisory functions, or first-line support – precisely the categories most amenable to AI substitution – this is not a theoretical risk but an imminent commercial reality.

The best board practice, in his view, is to begin that conversation now, before the disruption arrives. Boards of companies with largely physical revenue streams (manufacturing, infrastructure, logistics) have more time but cannot assume immunity: the service and support layers that wrap around physical products are already at risk. The companies he observed to be best prepared were those whose boards had made the disruption of their own top line a standing agenda item – asking

not 'are we deploying AI?' but 'are we disrupting ourselves, or are we waiting to be disrupted?'

4.4 Process Redesign: The Kaizen Parallel

One of the most practically grounded contributions came from the representative of a global industrial manufacturer with a century-old Kaizen heritage. The company – the world's largest producer of plasterboard and second largest in insulation materials – began its lean management journey in the 1980s, after a visit to a Japanese supplier in which an engineer's puzzlement at a question about process waste revealed a fundamentally different mindset: 'What seam?' The story has been repeated so often within the company that it has become a cultural touchstone – an invitation to question the necessity of every step in every process.

That mindset, she argued, is precisely what AI adoption requires. The willingness to ask 'do I even need to do this this way?' – rather than simply adding AI tools onto existing processes – is the cultural precondition for genuine transformation. She also emphasised the importance of project management discipline: without structured workstreams, defined timelines, and clear cross-silo ownership, organisations end up with pockets of AI-enabled efficiency that cannot connect to each other, producing partial optimisation rather than systemic change.

"A bad agent will only amplify a bad process. Companies have deployed AI technology on top of all processes without changing how anyone works. That is where the value is trapped."

4.5 The Cybersecurity Pole: From Afterthought to Strategic Priority

The session devoted significant attention to cybersecurity – a dimension that several participants felt was still treated as an afterthought in most board-level AI discussions. One participant, who sits on the World Economic Forum's AI Council and works closely with technology and cyber companies, described a concerning pattern: boards, driven by FOMO (fear of missing out) or fear of obsolescence, push management to deploy AI quickly; management responds by opening access broadly; CISOs are left to manage risks that are growing faster than their budgets and their authority.

The problem is structural as well as attitudinal. In many traditional enterprises, the CISO sits three or four layers below the board, with no meaningful channel to raise strategic concerns. The CISOs she meets at conferences describe feeling trapped: they cannot be the entity that blocks the business from moving, but they are acutely aware that every new AI deployment opens new attack surfaces. The threat

landscape is no longer just external; it is internal too. As more employees across an organisation gain access to AI tools, they interact directly with company data banks in ways that create data exposure risks that are difficult to audit and even harder to reverse.

The more advanced models – including frontier systems launched in mid-2026 – demonstrate capabilities that outpace the response time of traditional security teams. The shift required is from reactive to anticipatory: building security architectures that can absorb the expansion of AI access before, not after, the attack. Several participants noted that in regulated industries – financial services, utilities, telecoms – the EU's NIS2 Directive and related frameworks are already creating criminal liability for boards that fail to meet cybersecurity standards, providing an exogenous forcing mechanism that may accelerate board engagement with the topic.

A participant from a financial institution added a dimension specific to legacy industries: the challenge is not only building the right AI architecture but doing so on top of technology infrastructure that was never designed for AI integration. The data protection requirements, the system complexity, and the trust obligations of financial services clients create a layered set of constraints that slow deployment – but that also, if managed well, create a structural trust advantage over AI-native competitors.

4.6 Vendor Selection as Strategic Decision

The cybersecurity discussion led naturally to a broader point about vendor selection. In a landscape of 18,000 companies claiming AI capability, the choice of technology partner is itself a strategic and governance decision. Participants emphasised several dimensions of due diligence that are not yet standard board practice: understanding precisely how a vendor accesses and stores company data; assessing the accountability framework should an AI system hallucinate or be compromised; evaluating the vendor's own governance and ethics infrastructure; and considering the regulatory exposure of a vendor's home jurisdiction in the context of cross-border data flows.

One participant observed that commercial contracts have not kept pace with the liability questions AI raises. Boards should be asking their legal teams not only 'can we deploy this?' but 'who bears the risk when it goes wrong, and do our contracts reflect that allocation?'

Across these perspectives, a common pattern emerges: organisations are experimenting rapidly, but few have yet translated experimentation into coordinated, enterprise-wide transformation.

5. THE EARLY TALENT CRISIS: A LEADERSHIP GAP IN DISGUISE

One of the most direct exchanges of the session addressed a question that has become increasingly visible in public discourse: what happens to the pipeline of early talent when the apprenticeship model – the informal process by which junior employees learn by doing the routine work – is automated away?

The data is stark. In the US, approximately half of last year's graduating cohort had not secured employment at the time of the session. Youth unemployment rates are significantly above the national average. And when companies announce AI-driven headcount reductions, equity markets typically reward them – a signal that financial markets have already priced in labour displacement as a productivity gain.

Nathalie's response was characteristically frank. The displacement is real, but it is primarily a leadership gap, not an inevitable consequence of the technology. The implication is clear: workforce disruption is not predetermined by technology, but shaped by leadership choices.

At the company level, the leadership response has been deliberate. Three months before the session, the CEO made a specific commitment to double down on early-career hiring, launching what they call the Builder Programme. Rather than replicate the traditional apprenticeship model – in which junior employees learn by doing the routine, repetitive work that AI now handles – the programme assigns early-career talent to build product, run prompt engineering, and iterate on agent design. A cohort of 1,000 interns was arriving on campuses the week following the session for summer programmes structured entirely around building, not administering.

Reverse mentorship has also been formalised: junior employees are paired with senior professionals not simply to learn from them, but to teach them. The AI-native instincts of recent graduates are a resource for the organisation, not merely a nice-to-have.

Jean-Christophe added the historical perspective. In each of the three preceding industrial revolutions, decade over decade, more jobs were created than destroyed. Current Andreessen Horowitz and the Fed Reserve Bank of Atlanta data suggests that 91% of companies report no net job destruction attributable to AI, and that between 2% and 7% of jobs will ultimately be eliminated. The IMF and ILO suggest that between 25% and 40% of all jobs will change in structure significantly – and the responsibility for equipping people to navigate that change falls jointly on companies, learning institutions, and governments. The pace of reskilling is the critical variable.

6. THE MID-MARKET AND THE ECOSYSTEM CHALLENGE

Several participants raised a concern that is easy to overlook when the conversation is dominated by the experiences of large, well-resourced global companies: the position of the European mid-market. The backbone of the Swiss, German, and

broader continental European economy are the Mittelstand – mid-sized companies with revenues of €500 million to €1 billion, often niche global champions in their sectors, typically with limited technology infrastructure and constrained access to talent capable of driving AI transformation.

These companies face a set of compounding disadvantages: limited employer brand to attract AI talent, constrained resources to build internal infrastructure at scale, and insufficient size to fully leverage enterprise AI platforms designed for large-scale deployments. Yet they are not too small to be disrupted. Their competitive advantage is precision and depth of domain knowledge – precisely the qualities that AI is increasingly able to approximate.

Jean-Christophe argued that the response lies in partnership and in what he called organisational humility. The companies best positioned to survive and thrive in the AI transition are those whose leadership – starting with the CEO and the board – can acknowledge that being the world's best in a narrow niche is not, in itself, a sufficient response to AI disruption. The question is not whether to engage with AI, but with whom to build the capability. The Adecco Group's partnership with Salesforce – which has extended to building a joint venture – was offered as an example of what ecosystem collaboration looks like in practice: not simply buying software, but co-developing capability, sharing learning, and accessing agent infrastructure that would have been prohibitively expensive to build alone.

The same participant observed that one of the most persistent structural weaknesses of European corporate leadership is the reluctance to ask for help. Boards and CEOs regard dependence on external partners as a sign of weakness; the cultural default is self-sufficiency. In the current environment, that instinct is dangerous. The companies that move fastest are those most willing to learn, to partner, and to admit what they do not yet know.

7. BOARD GOVERNANCE IN THE AI ERA: WHAT NEEDS TO CHANGE

The final portion of the discussion focused specifically on the board itself. If management faces a transformation challenge, boards face a readiness challenge of a different kind: how do they acquire sufficient AI literacy to ask the right questions, set meaningful strategy, and provide genuine oversight – without becoming operational? The central question is no longer whether boards should engage with AI, but whether they are structurally equipped to do so effectively.

7.1 Recruiting for AI Fluency

Several participants pointed to recruitment as the single most powerful lever available to a board. Boards are constituted through deliberate choices about who sits at the table. Companies that have brought younger directors – including

executives with Silicon Valley exposure or AI-native career trajectories – onto their boards report that the quality of the AI conversation at board level changes qualitatively. One European company cited the appointment of a director under 40 with a Silicon Valley background as having transformed the energy of the AI discussion in the boardroom. The appointment was also, as an aside, noted to have lowered the average age of the board in a way that was itself productive.

The challenge is that board composition changes slowly, and the pace of AI evolution does not accommodate slow change. Several participants recommended a complementary approach: bringing management – particularly those who are daily users of AI – into the boardroom to present to directors on what they are experiencing. Bottom-up learning through management exposure may be a faster route to board AI fluency than waiting for the next board renewal cycle.

7.2 Establishing Dedicated Oversight Structures

One participant, a member of the supervisory board of a major industrial company, described the creation of a dedicated AI and Digital committee at board level – a structure she had not seen replicated in many comparable companies. The committee provides a focused channel for AI-related reporting and decision-making and has been accompanied by the intentional recruitment of directors with AI expertise into the board itself. She described both as the foundation of the foundation of the foundation: without specific attention and specific competence at board level, the oversight function in this domain is nominal rather than real.

A NASDAQ-listed technology company was cited as a broader reference point. The discussion referred not to an internal initiative of the company itself, but rather to the Governance Centre of Excellence available via the BoardVantage platform, which many boards use. This resource provides a collection of governance-related learning materials intended to support board members in adapting their practices in the AI era. The traditional board meeting format – dense pre-read materials, lengthy executive presentations, minimal time for questions – is, in that context, increasingly seen as outdated. AI can generate question frameworks from board papers, synthesise management reporting, and flag material risks in advance of a meeting. A board that continues to operate as if in a pre-AI environment risks inefficiency and, more importantly, may governing a fundamentally different organisation from the one that exists.

7.3 Reporting Metrics That Matter

Nathalie outlined the five reporting buckets through which her organisation gives its board meaningful visibility on the human capital dimensions of AI transformation. These are: technology adoption (daily AI usage rates across the workforce); work redesign (jobs redesigned end-to-end, decisions automated); workforce impact

(capacity unlocked, internal mobility rates, reskilling volumes); economic outcome (revenue per employee); and ethics, risk and accountability (aligned with frameworks such as the NIST AI RMF). The specificity of this reporting framework allows the board to engage with substance rather than with aspiration. This emphasis on measurable adoption, impact and accountability reflects the principles outlined in Salesforce's AI Fluency Playbook, which stresses that effective human-AI collaboration must be both observable and systematically tracked continuously optimised.

Several participants endorsed the principle of framing board questions at the task level rather than the technology level. Asking 'what AI have you deployed?' produces an inventory. Asking 'what work has been fundamentally redesigned through AI?' produces accountability. The quality of the board's question determines the quality of the management's answer.

8. CLOSING REFLECTIONS: TRUST, OPTIMISM AND THE WORK AHEAD

The session closed with reflections from both co-moderators on the themes that had emerged across ninety minutes of discussion.

Jean-Christophe returned to the imperative of human centricity. The companies most likely to navigate the AI transition successfully, he argued, are not those that move fastest, but those that move most intentionally – continuously developing the competencies that the transition requires, measuring impact rigorously, and maintaining empathy for the human experience of disruption. Running the transformation like a machine, without genuine attention to how people grow into new roles and new identities, is a path to failure. Ethical committees, he added, are not optional extras: every AI use case should pass through an ethics filter before it is scaled.

Nathalie closed on the theme of trust. Disruption is not a future event; it is the current condition. The question is how to manage it for the greater good. She expressed qualified optimism: in her direct experience, new jobs are already emerging that did not exist six months ago. A net increase in employment is, in her view, the likely long-run outcome – consistent with every prior wave of industrial transformation. But the path there is disruptive, and navigating it requires clear guiding principles, honest communication about risk appetite, and an unwavering commitment to bringing people along.

She was asked directly about Salesforce's own headcount trajectory. The answer was instructive: in two years of intensive AI transformation, total headcount is slightly up. Some functions – sales, forward deployed engineering – are substantially larger than they were. The company has not become smaller through AI; it has become differently configured. That reconfiguration is the work that boards need to mandate,

track, and champion – not as a cost exercise, but as the central strategic challenge of the decade.

"Disruption is happening everywhere right now. But how do you mitigate the amount of disruption for the greater good? Bringing people along is the most important thing."

9. SUMMARY OF KEY INSIGHTS

The following themes distil the principal insights from the session:

- Redesign is the unlock. Deploying AI on top of existing processes captures a fraction of the available value. The transformational opportunity lies in fundamental job and process redesign – asking what each role would look like if built from scratch today, knowing what we now know about AI. Companies that have made this shift report measurable gains in productivity, capacity and commercial performance.
- The 4Rs framework provides a replicable framework. Redesign, Reskill, Redeploy, Rebalance: this sequence offers boards and management teams a structured way to govern the workforce transition at scale, with clear reporting metrics at each stage.
- Governance must be both enabling and protective. The Ethical and Humane Use Council model – asking 'just because you can, should you?' – provides a template for boards to discharge their AI oversight obligations without becoming operational. The key discipline is maintaining the human in the loop at points of consequential decision-making.
- Cybersecurity is a strategic, not a technical, priority. As AI access broadens across organisations, the attack surface expands and the threat evolves faster than traditional security models can absorb. Boards in every sector – not only regulated industries – need to treat cybersecurity as a first-order governance concern, with the CISO appropriately elevated and resourced.
- The early talent pipeline requires deliberate intervention. The automation of routine apprenticeship-level work is real and rapid. Companies that are succeeding are those that have redesigned the early-career experience rather than simply reduced early-career hiring – creating new entry points that are substantive, developmental and genuinely valued.
- The mid-market needs an ecosystem response. Companies that lack the scale to build AI capability alone need to build it together – through strategic partnerships, co-development models, and ecosystem collaboration.

Organisational humility, beginning with the board and CEO, is a competitive advantage in the current environment.

- Board composition and literacy are lagging indicators of readiness. The surest path to meaningful AI governance at board level is through two parallel tracks: recruiting directors with genuine AI fluency and structuring regular bottom-up learning from management who are living the transformation daily.
- Purpose is not a soft variable. In conditions of rapid technological change and rising workforce uncertainty, organisational purpose is a measurable driver of retention, engagement and performance. The companies that manage the transition best will be those that maintain a coherent sense of why – a narrative that holds across every phase of the maturity curve.
- AI fluency is becoming a governance capability. Boards that cannot understand, question and measure human-AI collaboration risk losing meaningful oversight of organisational performance.

Acknowledgements

This knowledge paper synthesises insights from the Board Agora discussion session held on 15 June 2026. We extend our gratitude to Nathalie Scardino and Jean-Christophe Deslarzes for their expert contributions, and to all participating board members whose candid reflections and challenges enriched this dialogue. The quality of a Chatham House conversation depends entirely on the willingness of participants to speak with genuine honesty.

Board Agora Knowledge Paper

*AI's Impact on Talent and the Future of Work | June 2026 | Operating under
Chatham House Rules*